

PROYECTO CIENCIA DE DATOS - Sistema de Asistencia para Limitados Visuales en la Unab

Integrantes:

- Angiela Vargas Pérez
- Juan Pérez Bayona

Docente:

- Alfredo Díaz Claro

DESCRIPCIÓN:

El objetivo es desarrollar una aplicación de ayuda para personas con discapacidad visual en la Universidad (Unab), que permita conocer de forma interactiva el Laboratorio Smart Region Lab. La aplicación capturará imágenes de diferentes áreas, identificará espacios y equipos, y generará información hablada acerca de cada uno. Además, contará con un sistema interactivo basado en una red generativa de voz y texto que responderá preguntas y ofrecerá más detalles a los usuarios.

OBJETIVOS:

Accesibilidad y Orientación:

Permitir que personas con limitaciones visuales puedan conocer los distintos espacios, áreas y equipos del Smart Region Lab mediante descripciones de audio.

Captura y Procesamiento de Imágenes:

Utilizar técnicas de visión por computadora para reconocer y clasificar áreas y equipos a partir de fotografías.

Generación de Información:

Integrar un sistema de generación de voz y texto (mediante modelos generativos) que no solo describa los espacios, sino que también responda preguntas y amplíe la información según la interacción del usuario.

Interactividad y Respuesta Dinámica:

Ofrecer un sistema interactivo en el que el usuario pueda realizar consultas y recibir respuestas detalladas sobre los equipos, áreas y funciones de cada espacio.

✓ EXPLORACIÓN Y PREPROCESAMIENTO DE LA DATASET

Se tomaron muestras de cada clase, son 20 clases, las cuales se aumentaron, luego se proporcionaron con un tamaño de imagen de (224,224). Posterior a esto se realizó labelling a partir de bounding boxes mediante la plataforma gratuita de Roboflow, en la cual se asignaron los datos de entrenamiento, de validación y de prueba.

999 Total Images

[View All Images →](#)



Dataset Split

TRAIN SET

80%

798 Images

VALID SET

16%

159 Images

TEST SET

4%

42 Images

✓ Importar librerías

```
!pip install ultralytics opencv-python
```



Collecting ultralytics

Downloading ultralytics-8.3.133-py3-none-any.whl.metadata (37 kB)

Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)

```

Requirement already satisfied: numpy>=1.23.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (2.0.2)
Requirement already satisfied: matplotlib>=3.3.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (3.10.0)
Requirement already satisfied: pillow>=7.1.2 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (11.2.1)
Requirement already satisfied: pyyaml>=5.3.1 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (6.0.2)
Requirement already satisfied: requests>=2.23.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (2.32.3)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (1.15.3)
Requirement already satisfied: torch>=1.8.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (2.6.0+cu124)
Requirement already satisfied: torchvision>=0.9.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (0.21.0+cu124)
Requirement already satisfied: tqdm>=4.64.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (4.67.1)
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from ultralytics) (5.9.5)
Requirement already satisfied: py-cpuinfo in /usr/local/lib/python3.11/dist-packages (from ultralytics) (9.0.0)
Requirement already satisfied: pandas>=1.1.4 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (2.2.2)
Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.11/dist-packages (from ultralytics) (0.13.2)
Collecting ultralytics-thop>=2.0.0 (from ultralytics)
  Downloading ultralytics_thop-2.0.14-py3-none-any.whl.metadata (9.4 kB)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (4.53.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (1.4.6)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.3.0->ultralytics) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.1.4->ultralytics) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.1.4->ultralytics) (2025.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests>=2.23.0->ultralytics) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests>=2.23.0->ultralytics) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests>=2.23.0->ultralytics) (2.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests>=2.23.0->ultralytics) (2025.2.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (4.12.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (2025.3.2)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==12.3.1.170 (from torch>=1.8.0->ultralytics)
  Downloading nvidia_cusparselt_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 (from torch>=1.8.0->ultralytics)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=1.8.0->ultralytics)

```

```
%pip show ultralytics
```

```

Name: ultralytics
Version: 8.3.133
Summary: Ultralytics YOLO for SOTA object detection, multi-object tracking, instance segmentation, pose estimation and image classification
Home-page: https://ultralytics.com
Author:
Author-email: Glenn Jocher <glenn.jocher@ultralytics.com>, Jing Qiu <jing.qiu@ultralytics.com>
License: AGPL-3.0
Location: /usr/local/lib/python3.11/dist-packages
Requires: matplotlib, numpy, opencv-python, pandas, pillow, psutil, py-cpuinfo, pyyaml, requests, scipy, seaborn, torch, torchvision, tqdm
Required-by:

```

```

import torch
import os
import cv2
import shutil
import zipfile
import requests
import random
from ultralytics import YOLO
from IPython.display import Image, clear_output

```

```

Creating new Ultralytics Settings v0.6 file
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see https://docs.ultralytics.com/quickstart

```

✎ Importar la dataset

La dataset se encuentra organizada de la siguiente manera:


```
dataset/
├── train/
│   ├── images/
│   └── labels/
├── valid/
│   ├── images/
│   └── labels/
├── test/
│   ├── images/
│   └── labels/
├── data.yaml
├── README.dataset
└── README.roboflow
```

```
# Definir rutas base
base_path = "/content"
zip_path = os.path.join(base_path, "LABCD.v2i.yolov8.zip")

def download_zip(url, dest):
    response = requests.get(url, stream=True)
    with open(dest, "wb") as file:
        for chunk in response.iter_content(chunk_size=8192):
            file.write(chunk)
    print("Archivo ZIP descargado correctamente.")

def extract_zip(zip_file, dest_folder):
    dataset_folder = os.path.join(dest_folder, "dataset") # Extraer en "dataset/"
    os.makedirs(dataset_folder, exist_ok=True) # Asegurar que la carpeta existe
    with zipfile.ZipFile(zip_file, 'r') as zip_ref:
        zip_ref.extractall(dataset_folder) # Extraer en dataset/
    print(f"Contenido extraído en {dataset_folder}: {os.listdir(dataset_folder)}")


# Ejecutar procesos
download_zip("https://github.com/Angiela-Vargas/Data-Science-LABCD/blob/main/LABCD.v2i.yolov8.zip?raw=true", zip_path)
extract_zip(zip_path, base_path)
```

 Archivo ZIP descargado correctamente.
Contenido extraído en /content/dataset: ['data.yaml', 'valid', 'test', 'README.roboflow.txt', 'README.dataset.txt', 'train']

✎ Verificar la descarga y extracción

```
labels_val_path = "/content/dataset/valid/labels" # Ajusta la ruta según tu dataset

if os.path.exists(labels_val_path):
    label_files = [f for f in os.listdir(labels_val_path) if f.endswith(".txt")]
    if len(label_files) == 0:
        print("⚠️ No se encontraron etiquetas en:", labels_val_path)
    else:
        print(f"✅ Se encontraron {len(label_files)} archivos de etiquetas.")
else:
    print("⚠️ La carpeta de etiquetas no existe.")
```

 ✅ Se encontraron 159 archivos de etiquetas.

✎ Validar estructura completa del dataset

```

base_path = "/content/dataset"
subdirs = ["train/images", "train/labels", "valid/images", "valid/labels", "test/images", "test/labels"]

for subdir in subdirs:
    full_path = os.path.join(base_path, subdir)
    if not os.path.exists(full_path):
        print(f"🔥 ERROR: No se encontró la carpeta {full_path}")
    else:
        print(f"✅ {full_path} encontrado, contiene {len(os.listdir(full_path))} archivos")

```

```

✅ /content/dataset/train/images encontrado, contiene 798 archivos
✅ /content/dataset/train/labels encontrado, contiene 798 archivos
✅ /content/dataset/valid/images encontrado, contiene 159 archivos
✅ /content/dataset/valid/labels encontrado, contiene 159 archivos
✅ /content/dataset/test/images encontrado, contiene 42 archivos
✅ /content/dataset/test/labels encontrado, contiene 42 archivos

```

Edición del archivo dataset.yaml

El archivo .yaml define:

- Ruta base de los datos.
- Subcarpetas de entrenamiento, validación y prueba.
- Número de clases (nc).
- Nombres de las clases.

```
```yaml
```

```
train: ../train/images
```

```
val: ../valid/images
```

```
test: ../test/images
```

```
nc: 21
```

```
names: ['Microscopio óptico avanzado', 'Agitador magnético con placa calefactora', 'Aire acondicionado YORK', 'Cafetera', 'Computador',
'Control remoto de robot', 'Control remoto de Smart TV', 'Estación de soldadura por aire caliente', 'Impresora 3D Creality K1 Max', 'Mini
fresadora CNC', 'Mouse', 'Objetos (varios)', 'Osciloscopio digital Rigol DS1202ZE', 'Puerta', 'Robot humanoide AULER', 'Robot Lego Mindstorms
EV3', 'Router', 'Sensor de suelo integrado', 'Tablet', 'Tomacorrientes 110V', 'UPS Netio']
```

```
roboflow: workspace: labcd project: labcd version: 2 license: CC BY 4.0 url: https://universe.roboflow.com/labcd/labcd/dataset/2
```

## Nombres:

- 0: Microscopio óptico avanzado
- 1: Agitador magnético con placa calefactora
- 2: Aire acondicionado YORK
- 3: Cafetera
- 4: Computador
- 5: Control remoto de robot
- 6: Control remoto de Smart TV
- 7: Estación de soldadura por aire caliente
- 8: Impresora 3D Creality K1 Max
- 9: Mini fresadora CNC
- 10: Mouse
- 11: Objetos (varios)
- 12: Osciloscopio digital Rigol DS1202ZE
- 13: Puerta
- 14: Robot humanoide AULER
- 15: Robot Lego Mindstorms EV3
- 16: Router

17: Sensor de suelo integrado  
 18: Tablet  
 19: Tomacorrientes 110V  
 20: UPS Netio

## ✓ Cargar modelo base YOLOv8 y Cargar pesos preentrenados YOLOv8

Se carga el modelo YOLOv8 con pesos base (yolov8s.pt), listos para ser ajustados a nuestro conjunto de datos.

```
model = YOLO("yolov8s.pt")
```

Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8s.pt> to 'yolov8s.pt'...  
 100%|██████████| 21.5M/21.5M [00:00<00:00, 95.9MB/s]

## ✓ Entrenamiento del modelo YOLOv8 con la dataset creada en Roboflow

```
model.train(data="/content/dataset/data.yaml", epochs=50, imgsz=640, batch=16)
```

Ultralytics 8.3.133 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)  
**engine/trainer:** agnostic\_nms=False, amp=True, augment=False, auto\_augment=randaug, batch=16, bgr=0.0, box=7.5, cache=False, cfg=None  
 Downloading <https://ultralytics.com/assets/Arial.ttf> to '/root/.config/Ultralytics/Arial.ttf'...  
 100%|██████████| 755k/755k [00:00<00:00, 20.7MB/s]Overriding model.yaml nc=80 with nc=21

	from	n	params	module	arguments
0	-1	1	928	ultralytics.nn.modules.conv.Conv	[3, 32, 3, 2]
1	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
2	-1	1	29056	ultralytics.nn.modules.block.C2f	[64, 64, 1, True]
3	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
4	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
5	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
6	-1	2	788480	ultralytics.nn.modules.block.C2f	[256, 256, 2, True]
7	-1	1	1180672	ultralytics.nn.modules.conv.Conv	[256, 512, 3, 2]
8	-1	1	1838080	ultralytics.nn.modules.block.C2f	[512, 512, 1, True]
9	-1	1	656896	ultralytics.nn.modules.block.SPPF	[512, 512, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	591360	ultralytics.nn.modules.block.C2f	[768, 256, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
16	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
19	-1	1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]
21	-1	1	1969152	ultralytics.nn.modules.block.C2f	[768, 512, 1]
22	[15, 18, 21]	1	2124175	ultralytics.nn.modules.head.Detect	[21, [128, 256, 512]]

Model summary: 129 layers, 11,143,727 parameters, 11,143,711 gradients, 28.7 GFLOPs

Transferred 349/355 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

**AMP:** running Automatic Mixed Precision (AMP) checks...

Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolo11n.pt> to 'yolo11n.pt'...  
 100%|██████████| 5.35M/5.35M [00:00<00:00, 120MB/s]

**AMP:** checks passed ✓

**train:** Fast image access ✓ (ping: 0.0±0.0 ms, read: 350.8±108.0 MB/s, size: 9.7 KB)

**train:** Scanning /content/dataset/train/labels... 798 images, 0 backgrounds, 0 corrupt: 100%|██████████| 798/798 [00:00<00:00, 2430.51i

**augmentations:** Blur(p=0.01, blur\_limit=(3, 7)), MedianBlur(p=0.01, blur\_limit=(3, 7)), ToGray(p=0.01, method='weighted\_average', num\_val: Fast image access ✓ (ping: 0.0±0.0 ms, read: 247.0±129.8 MB/s, size: 7.9 KB)

**val:** Scanning /content/dataset/valid/labels... 159 images, 0 backgrounds, 0 corrupt: 100%|██████████| 159/159 [00:00<00:00, 1646.72it/

Plotting labels to runs/detect/train/labels.jpg...

**optimizer:** 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' aut  
**optimizer:** AdamW(lr=0.0004, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias(decay=0.0)

Image sizes 640 train, 640 val

Using 2 dataloader workers

Logging results to runs/detect/train

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/50	3.73G	1.054	4.52	1.552	41	640: 100% ██████████  50/50 [00:17<00:00, 2.89it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████  5/5 [00:02<00:00, 1.79it/s]

```

metrics = model.val()
print(metrics)

```

Ultralitics 8.3.133 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)  
 Model summary (fused): 72 layers, 11,133,711 parameters, 0 gradients, 28.5 GFLOPs  
 val: Fast image access (ping: 0.0±0.0 ms, read: 315.1±130.2 MB/s, size: 7.0 KB)  
 val: Scanning /content/dataset/valid/labels.cache... 159 images, 0 backgrounds, 0 corrupt: 100%|██████████| 159/159 [00:00<?, ?it/s]

	Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)
	all	159	194	0.987	0.957	0.982	0.906
advanced_optical_microscope	9	9	0.991	1	0.995	0.89	
agitador_magnetico_placa_calentamiento_EQ_AMPC1C	9	9	0.991	1	0.995	0.884	
aire_acondicionado_YORK	6	6	0.984	1	0.995	0.96	
cafetera	8	8	0.994	1	0.995	0.949	
computador	10	19	0.944	0.895	0.983	0.827	
control_remoto_robot	11	11	0.988	0.727	0.79	0.732	
control_remoto_smart_tv	8	8	0.994	1	0.995	0.918	
estacion_soldadura_aire_caliente	10	10	0.992	1	0.995	0.909	
impresora_3d_creality_k1_max	8	8	1	0.884	0.995	0.841	
mini_fresadora_cnc	8	8	0.989	1	0.995	0.885	
mouse	8	8	0.988	1	0.995	0.956	
objects	1	1	0.922	1	0.995	0.995	
osciloscopio_digital_rigol_ds1202ze	7	7	1	0.884	0.995	0.899	
puerta	8	8	0.991	1	0.995	0.828	
robot_humanoide_AULER	14	14	1	0.87	0.935	0.85	
robot_lego_mindstorms_ev3	8	8	0.991	1	0.995	0.862	
router	8	8	0.99	1	0.995	0.946	
soil_integrated_sensor	8	8	0.988	1	0.995	0.975	
tablet	8	8	0.985	1	0.995	0.981	
tomacorrientes_110V	8	20	1	0.943	0.995	0.957	
ups_netio	8	8	1	0.896	0.995	0.976	

Speed: 3.3ms preprocess, 9.8ms inference, 0.0ms loss, 2.9ms postprocess per image  
 Results saved to runs/detect/train2  
 ultralytics.utils.metrics.DetMetrics object with attributes:

ap\_class\_index: array([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20])  
 box: ultralytics.utils.metrics.Metric object  
 confusion\_matrix: <ultralitics.utils.metrics.ConfusionMatrix object at 0x795891970710>  
 curves: ['Precision-Recall(B)', 'F1-Confidence(B)', 'Precision-Confidence(B)', 'Recall-Confidence(B)']  
 curves\_results: [[array([
 0, 0.001001, 0.002002, 0.003003, 0.004004, 0.005005, 0.006006, 0.007007, 0.
 0.024024, 0.025025, 0.026026, 0.027027, 0.028028, 0.029029, 0.03003, 0.031031, 0.032032, 0.03303
 0.048048, 0.049049, 0.05005, 0.051051, 0.052052, 0.053053, 0.054054, 0.055055, 0.056056, 0.05705
 0.072072, 0.073073, 0.074074, 0.075075, 0.076076, 0.077077, 0.078078, 0.079079, 0.08008, 0.08108
 0.096096, 0.097097, 0.098098, 0.099099, 0.1001, 0.1011, 0.1021, 0.1031, 0.1041, 0.1051
 0.12012, 0.12112, 0.12212, 0.12312, 0.12412, 0.12513, 0.12613, 0.12713, 0.12813, 0.1291
 0.14414, 0.14515, 0.14615, 0.14715, 0.14815, 0.14915, 0.15015, 0.15115, 0.15215, 0.1531
 0.16817, 0.16917, 0.17017, 0.17117, 0.17217, 0.17317, 0.17417, 0.17518, 0.17618, 0.1771
 0.19219, 0.19319, 0.19419, 0.1952, 0.1962, 0.1972, 0.1982, 0.1992, 0.2002, 0.201
 0.21622, 0.21722, 0.21822, 0.21922, 0.22022, 0.22122, 0.22222, 0.22322, 0.22422, 0.2252
 0.24024, 0.24124, 0.24224, 0.24324, 0.24424, 0.24525, 0.24625, 0.24725, 0.24825, 0.2492
 0.26426, 0.26527, 0.26627, 0.26727, 0.26827, 0.26927, 0.27027, 0.27127, 0.27227, 0.2732
 0.28829, 0.28929, 0.29029, 0.29129, 0.29229, 0.29329, 0.29429, 0.2953, 0.2963, 0.297
 0.31231, 0.31331, 0.31431, 0.31532, 0.31632, 0.31732, 0.31832, 0.31932, 0.32032, 0.3213
 0.33634, 0.33734, 0.33834, 0.33934, 0.34034, 0.34134, 0.34234, 0.34334, 0.34434, 0.3453
 0.36036, 0.36136, 0.36236, 0.36336, 0.36436, 0.36537, 0.36637, 0.36737, 0.36837, 0.3693
 0.38438, 0.38539, 0.38639, 0.38739, 0.38839, 0.38939, 0.39039, 0.39139, 0.39239, 0.3933
 0.40841, 0.40941, 0.41041, 0.41141, 0.41241, 0.41341, 0.41441, 0.41542, 0.41642, 0.4174
 0.43243, 0.43343, 0.43443, 0.43544, 0.43644, 0.43744, 0.43844, 0.43944, 0.44044, 0.4414
 0.45646, 0.45746, 0.45846, 0.45946, 0.46046, 0.46146, 0.46246, 0.46346, 0.46446, 0.4654
 0.48048, 0.48148, 0.48248, 0.48348, 0.48448, 0.48549, 0.48649, 0.48749, 0.48849, 0.4894
 0.5045, 0.50551, 0.50651, 0.50751, 0.50851, 0.50951, 0.51051, 0.51151, 0.51251, 0.5135
 ])]

## ✓ Testear el modelo

```

Cargar la imagen
image_path = "/content/dataset/test/images/augmented_0_jpg.rf.bfdcb4435da150ef47c97b4b1c2ab783.jpg"

image = cv2.imread(image_path)
results = model.predict(image_path, save=True, imgsz=640)

Guardar la imagen
cv2.imwrite("imagen.jpg", image)

```

image 1/1 /content/dataset/test/images/augmented\_0\_jpg.rf.bfdcb4435da150ef47c97b4b1c2ab783.jpg: 640x640 1 puerta, 16.0ms  
 Speed: 4.8ms preprocess, 16.0ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train3  
 True

```
results[0].show()
```



Realizaremos una predicción de todas las imágenes que se encuentran en la carpeta de test

```
#Hagamos la predicción de todas las imágenes que están en test

Definir la ruta a la carpeta de imágenes de prueba
test_images_dir = "/content/dataset/test/images"

Iterar sobre todas las imágenes en la carpeta de prueba
for filename in os.listdir(test_images_dir):
 if filename.endswith((''.jpg', '.jpeg', '.png')): # Ajusta las extensiones si es necesario
 image_path = os.path.join(test_images_dir, filename)
 results = model.predict(image_path, save=True, imgsz=640) # Realizar la predicción
 print(f"Predicción realizada para: {filename}")
 # Puedes acceder a las predicciones a través de results[0].boxes, results[0].probs, etc.
 # Si deseas visualizar las imágenes con las predicciones:
 results[0].show()
```



image 1/1 /content/dataset/test/images/original\_15\_jpg.rf.acb378137bde5d264a582f7e5458adb.jpg: 640x640 1 puerta, 16.0ms  
 Speed: 2.7ms preprocess, 16.0ms inference, 1.7ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train4  
 Predicción realizada para: original\_15\_jpg.rf.acb378137bde5d264a582f7e5458adb.jpg



image 1/1 /content/dataset/test/images/original\_4\_jpg.rf.e09549b155ad6a60c032e23436495d22.jpg: 640x640 3 tomacorrientes\_110Vs, 16.0ms  
 Speed: 3.5ms preprocess, 16.0ms inference, 1.7ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train5  
 Predicción realizada para: original\_4\_jpg.rf.e09549b155ad6a60c032e23436495d22.jpg



image 1/1 /content/dataset/test/images/augmented\_21\_jpg.rf.f36c7adf607650bca9f3a9842fc56337.jpg: 640x640 1 advanced\_optical\_microscope, 16.0ms  
 Speed: 5.0ms preprocess, 16.0ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train6  
 Predicción realizada para: augmented\_21\_jpg.rf.f36c7adf607650bca9f3a9842fc56337.jpg




image 1/1 /content/dataset/test/images/original\_27\_jpg.rf.61885bf47ed65ba4e7840b289cc7aab0.jpg: 640x640 1 mini\_fresadora\_cnc, 16.0ms  
 Speed: 4.7ms preprocess, 16.0ms inference, 2.0ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train7  
 Predicción realizada para: original\_27\_jpg.rf.61885bf47ed65ba4e7840b289cc7aab0.jpg




image 1/1 /content/dataset/test/images/original\_5\_jpg.rf.f0dbdc2b69ee008d6ed8f30a5a5efe62.jpg: 640x640 1 tablet, 16.0ms  
 Speed: 4.1ms preprocess, 16.0ms inference, 2.0ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train8  
 Predicción realizada para: original\_5\_jpg.rf.f0dbdc2b69ee008d6ed8f30a5a5efe62.jpg







image 1/1 /content/dataset/test/images/original\_22\_jpg.rf.26d8c2b9f4f761749f26624f179c4b86.jpg: 640x640 1 aire\_acondicionado\_YORK, 15.7ms  
 Speed: 3.7ms preprocess, 15.7ms inference, 1.9ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train9  
 Predicción realizada para: original\_22\_jpg.rf.26d8c2b9f4f761749f26624f179c4b86.jpg

dicionado\_YORK 0.94



image 1/1 /content/dataset/test/images/original\_12\_jpg.rf.90a8cac4ca5a41bdf47e7869ee2556b3.jpg: 640x640 1 estacion\_soldadura\_aire\_caliente, 15.4ms  
 Speed: 3.6ms preprocess, 15.4ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train10  
 Predicción realizada para: original\_12\_jpg.rf.90a8cac4ca5a41bdf47e7869ee2556b3.jpg

ra\_aire\_caliente 0.91



image 1/1 /content/dataset/test/images/original\_9\_jpg.rf.252e00adb381960466ba9339f30a8a16.jpg: 640x640 1 soil\_integrated\_sensor, 15.4ms  
 Speed: 3.4ms preprocess, 15.4ms inference, 2.0ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train11  
 Predicción realizada para: original\_9\_jpg.rf.252e00adb381960466ba9339f30a8a16.jpg

tegrated\_sensor 0.95



image 1/1 /content/dataset/test/images/original\_8\_jpg.rf.0f3fc8cce4e1b4eddec95b1b3c2419d6.jpg: 640x640 1 soil\_integrated\_sensor, 15.4ms  
 Speed: 3.8ms preprocess, 15.4ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train12  
 Predicción realizada para: original\_8\_jpg.rf.0f3fc8cce4e1b4eddec95b1b3c2419d6.jpg

tegrated\_sensor 0.93



image 1/1 /content/dataset/test/images/original\_4\_jpg.rf.5944ee2318677ec8148813d88ecfb71f.jpg: 640x640 1 impresora\_3d\_creativity\_k1\_max, 15.4ms  
 Speed: 3.7ms preprocess, 15.4ms inference, 2.4ms postprocess per image at shape (1, 3, 640, 640)  
 Results saved to runs/detect/train13  
 Predicción realizada para: original\_4\_jpg.rf.5944ee2318677ec8148813d88ecfb71f.jpg





image 1/1 /content/dataset/test/images/original\_12\_jpg.rf.41011768ee60b69e3c8a3343fbe03ed9.jpg: 640x640 2 computadores, 15.5ms  
Speed: 3.5ms preprocess, 15.5ms inference, 1.9ms postprocess per image at shape (1, 3, 640, 640)

Results saved to runs/detect/train14

Predicción realizada para: original\_12\_jpg.rf.41011768ee60b69e3c8a3343fbe03ed9.jpg



image 1/1 /content/dataset/test/images/original\_25\_jpg.rf.f2c1b329cf9e893f5f105fe692c32485.jpg: 640x640 1 control\_remoto\_robot, 15.4ms  
Speed: 3.4ms preprocess, 15.4ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)

Results saved to runs/detect/train15

Predicción realizada para: original\_25\_jpg.rf.f2c1b329cf9e893f5f105fe692c32485.jpg



image 1/1 /content/dataset/test/images/augmented\_23\_jpg.rf.5947bc3c7fd153c9b407187fa0628f9c.jpg: 640x640 1 impresora\_3d\_crealty\_k1\_max, 15.4ms  
Speed: 3.1ms preprocess, 18.3ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)

Results saved to runs/detect/train16

Predicción realizada para: augmented\_23\_jpg.rf.5947bc3c7fd153c9b407187fa0628f9c.jpg




image 1/1 /content/dataset/test/images/original\_19\_jpg.rf.0e1ac3dc2949887e8fb82ba16d75935e.jpg: 640x640 1 router, 15.4ms  
Speed: 3.5ms preprocess, 15.4ms inference, 1.9ms postprocess per image at shape (1, 3, 640, 640)

Results saved to runs/detect/train17

Predicción realizada para: original\_19\_jpg.rf.0e1ac3dc2949887e8fb82ba16d75935e.jpg



image 1/1 /content/dataset/test/images/original\_16\_jpg.rf.1e719bdba9a24619a7508f67b5745c68.jpg: 640x640 1 robot\_humanoide\_AULER, 1 tomac  
Speed: 3.5ms preprocess, 15.4ms inference, 1.8ms postprocess per image at shape (1, 3, 640, 640)